

1 Probability and random variables

- **Probability:** S sample space (all possible states of the system), $\mathcal{F} \subset \mathcal{P}(S)$ a σ -algebra, $P : \mathcal{F} \rightarrow \mathbb{R}$ a measure, such that $P(S) = 1$.
- **Random variable:** $X : S \rightarrow \mathbb{R}$, such that preimages of open sets are in \mathcal{F} (i.e. has a well defined probability).
- **Cumulative distribution function** of random variable: $F_X(t) = P(X \leq t)$.
- **Probability distribution** of random variable: g such that $F_X(t) = \sum_{x \leq t, x \in C} g(x)$.
- **Probability density function:** f such that $F_X(t) = \int_{-\infty}^t f(s)ds$.
- Two random variables have the **same distribution** if they have the same cdf.

Example: **uniform distribution:**

- S a finite interval $[a, b]$
- \mathcal{F} : Set of Borel sets on S (sets with a well defined “length”)
- P : Borel measure (“length”) divided by $b - a$
- $X = id$.

1.1 Expectation of random variables and their functions

- X is a random variable, the **expectation** of X is $E[X] = \int_S X dP$.
- The **variance** of X is $E[(X - E[X])^2]$.
- The k -th **moment** of X is $E[X^k]$.
- The **moment generating function** of X is $E[e^{Xt}]$ (two sided Laplace transform)
- The **characteristic function** of X is $E[e^{itX}]$ (Fourier transform)

Since expectation is defined via integration, one can use the properties of integration to prove statements regarding expectation.

Example: **Chebyshev’s theorem:** $E[X] = 0$, $E[X^2] = 1$, then $P(|X| < k) \geq 1 - \frac{1}{k^2}$.
Proof:

$$1 = E[X^2] = \int_S X^2 dP \geq k^2 \int_{|X| \geq k} 1 dP = k^2(1 - P(|X| < k))$$

Example: If X has p.d.f. f_X , then $E[g(X)] = \int_{-\infty}^{\infty} g f_X dt$. We prove it when $g(X)$ is bounded via Fubini's theorem:

$$\begin{aligned} E[g(X)] &= \int_S g(X) dP \\ &= \int_{g(X) \geq 0} \int_0^{g(X)} 1 dy dP - \int_{g(X) < 0} \int_{g(X)}^0 1 dy dP \\ &= \int_0^{\infty} \int_{g^{-1}([y, \infty))} f_X(t) dt dy - \int_{-\infty}^0 \int_{g^{-1}([-\infty, y])} f_X(t) dt dy \\ &= \int_{-\infty}^{\infty} g f_X dt \end{aligned}$$

There is a multivariate version of this formula, and one can also write down $E[g(X)]$ when only the c.d.f. of X is known (via Fubini's theorem or integration by parts).

Can you write down a random variable with neither probability distribution nor p.d.f.?

Can you write down a random variable with no expectation?

1.2 Independence and conditional probability for random events

- $A, B \in \mathcal{F}$ are **independent** iff $P(A \cap B) = P(A)P(B)$.
- If $P(B) \neq 0$, $P(A \cap B) = P(B)P(A|B)$. Here $P(A|B)$ is the **conditional probability** of A when B is known to happen.

1.3 Joint distribution, marginal distribution, conditional distribution

1.3.1 Joint distribution

- X and Y are two random variables. The **joint cumulative distribution function** is $F(s, t) = P(X \leq s, Y \leq t)$.
- If $F(s, t) = \sum_{(x, y) \in C, x \leq s, y \leq t} g(s, t)$, we call g the **joint probability distribution**.
- If $F(s, t) = \int_{(-\infty, s] \times (-\infty, t]} f(x, y) dx dy$ we call f the **joint probability density function**.
- X and Y are called independent iff the joint c.d.f. is $F(x, y) = F_X(x)F_Y(y)$.
- The **covariance** between X and Y is $E[(X - E[X])(Y - E[Y])]$

Example: X and Y are two independent random variable with uniform distribution on $[0, 1]$. What is the joint distribution function of X and Y ? How about $\max(X, Y)$ and $\min(X, Y)$? What are their covariances?

1.3.2 Marginal distribution

Knowing the joint c.d.f. of X and Y , the c.d.f. of X or Y are called the **marginal cumulative distribution function**, their p.d. or p.d.f. the **marginal p.d. or marginal p.d.f.**

1.3.3 Conditional distribution

- If A is a set such that $P(Y \in A) > 0$, then the **conditional cumulative distribution function** of X is $F_{X|Y \in A}(t) = P(X \leq t | Y \in A) = P(X \leq t \cap Y \in A) / P(Y \in A)$. The **conditional p.d.f.**, **conditional p.d.** and **conditional expectation** are defined similarly.
- If $P(Y \in A) = 0$ there isn't a definition of conditional distribution that works in all cases. For example, if X, Y has joint p.d.f. $f_{X,Y}$, and the marginal p.d.f. of Y , denoted as $f_Y(y) = \int_{\mathbb{R}} f_{X,Y}(x, y) dx$, exists and is non zero at y_0 , then the conditional p.d.f. at $Y = y_0$ is defined as $f_{X|Y=y_0} = f_{X,Y}(x, y_0) / f_Y(y_0)$. The conditional c.d.f. is its integral.

Remark: The definition of conditional distribution for the case $P(Y \in A) = 0$ depends on Y and not just $Y^{-1}(A)$. For example, if $Z = Ye^X$, $f_{X|Y=0} \neq f_{X|Z=0}$.

Example: X is a random variable with uniform distribution on $[0, 1]$, $P(Y = 1 | X = p) = p$ (i.e. $P(Y = 1 | X \in A) = \int_A p dF_x(p)$), $P(Y = 0 | X = p) = 1 - p$. Find the conditional distribution of X when $Y = 1$.

When there are N random variables, $N \geq 3$, the joint/marginal/conditional distributions can be defined analogously.

2 Special probability distributions, central limit theorem

2.1 Special discrete distributions

- **Bernoulli distribution:** $f(1) = \theta$, $f(0) = 1 - \theta$.
- **Binomial distribution** (sum of iid Bernoulli): $f(x) = \binom{n}{x} \theta^x (1 - \theta)^{n-x}$, $x = 0, 1, \dots, n$.

- **Negative Binomial distribution** (waiting time for the k -th success of iid trials): $f(x) = \binom{x-1}{k-1} \theta^k (1-\theta)^{x-k}$, $x = k, k+1, \dots$. When $k = 1$ it is the **geometric distribution**.
- **Hypergeometric distribution** (randomly pick n elements at random from N elements, the number of elements picked from a fixed subset of M elements) $f(x) = \binom{M}{x} \binom{N-M}{n-x} \binom{N}{n}^{-1}$.
- **Poisson distribution** (limit of binomial as $n \rightarrow \infty$, $n\theta \rightarrow \lambda$) $f(x) = \lambda^x e^{-\lambda} / x!$.
- **Multinomial distribution** $f(x_1, \dots, x_k) = \binom{n}{x_1, \dots, x_k} \theta_1^{x_1} \dots \theta_k^{x_k}$, $\sum_i x_i = n$, $\theta_i \theta_i = 1$.
- **Multivariate Hypergeometric distribution** $f(x_1, \dots, x_k) = \prod_i \binom{M_i}{x_i} \binom{N}{n}^{-1}$. $\sum_i x_i = n$, $\sum_i M_i = N$.

2.2 Special continuous distributions

- **Uniform distribution**: $f(x) = \begin{cases} 1/(b-a) & x \in (a, b) \\ 0 & x \notin (a, b) \end{cases}$.
- **Normal distribution**: $f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$.
- **Multivariate Normal distribution**: $x \in \mathbb{R}^d$, Σ positive definite $d \times d$ symmetric matrix, $f(x) = (2\pi)^{-d/2} |\Sigma|^{-1/2} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1} (x-\mu)}$.
- **χ^2 distribution** d : degrees of freedom. Squared sum of d normal distributions: $f(x) = \begin{cases} \frac{1}{2^{d/2} \Gamma(d/2)} x^{\frac{d-2}{2}} e^{-x/2} & x > 0 \\ 0 & x \leq 0 \end{cases}$.
- **Exponential distribution** $f(x) = \begin{cases} \frac{1}{\theta} e^{-x/\theta} & x > 0 \\ 0 & x \leq 0 \end{cases}$.
- **Gamma-distribution**: $f(x) = \begin{cases} \frac{1}{\beta^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\beta} & x > 0 \\ 0 & x \leq 0 \end{cases}$.
- **Beta distribution**: (conjugate prior of Bernoulli distribution) $f(x) = \begin{cases} \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1} & x \in (0, 1) \\ 0 & x \notin (0, 1) \end{cases}$.

2.3 Law of Large Numbers and Central Limit Theorem

2.3.1 Convergence

- **Convergence in distribution:** cdf pointwise convergence.
- **Convergence almost surely:** $P(\lim_i X_i \neq X) = 0$.

Example: X uniform on $[0, 1]$, $Y_i = \begin{cases} 1 & \exists n \in \mathbb{Z} (X + n \in [\sum_{j=1}^i \frac{1}{j}, \sum_{j=1}^{i+1} \frac{1}{j}]) \\ 0 & \text{otherwise} \end{cases}$.

Then Y_i converges to 0 in distribution but not almost surely.

2.3.2 CLT and weak LLN

Levy's continuity theorem: If $\phi_{X_j} \rightarrow \phi_X$ pointwise, then X_j converges to X in distribution.

Weak Law of Large Numbers X_i i.i.d. with expectation μ . $S_n = \frac{1}{n} \sum_{i=1}^n X_i$. Then S_n converges to μ in distribution.

(Levy's) Central Limit Theorem X_i i.i.d. with expectation μ and variance $\sigma^2 > 0$. $Y_n = \sqrt{\frac{1}{n\sigma^2}} \sum_i (X_i - \mu)$, then Y_n converges in distribution to standard normal distribution (normal distribution with $\mu = 0$ and $\sigma^2 = 1$).

Proof of both theorems (assume X_i bounded): Taylor expansion of the characteristic function.

One can also use the continuity of moment generating function, which is the argument in the textbook.

2.3.3 Strong Law of Large Numbers

Borel-Cantelli Lemma A_i events, $i = 1, 2, \dots$, $\sum_i (A_i) < \infty$, then $P(\cap_i (\cup_{j>i} A_j)) = 0$. (the probability of infinitely many A_i happening is 0)

Proof: $P(\cap_i (\cup_{j>i} A_j)) \leq P(\cup_{j>i} A_j) \leq \sum_{j>i} P(A_j)$ which converges to 0 as $i \rightarrow \infty$.

Strong Law of Large Numbers $X_i, i = 1, 2, \dots$ i.i.d. (independent with identical distribution) and $E(X_i) = \mu$, then $Y_n = \frac{1}{n} \sum_{i=1}^n X_i$ converges a.s. to constant μ .

Proof (assume X_i bounded by M): Suppose $Var(X_i) = m$. $\sqrt{\frac{n}{m}}(Y_n - \mu)$ has expectation 0 and variance 1, so $P(|Y_n - \mu| > C\sqrt{\frac{m}{n}}) < 1/C^2$ by Chebyshev's theorem. Now let $n_k = k^4$, $C_k = k$, then $Y_{n_k} = Y_{k^4}$ converges a.s. to μ by Borel-Cantelli.

$Y_n = (\lfloor n^{1/4} \rfloor^4 Y_{\lfloor n^{1/4} \rfloor^4} + X_{\lfloor n^{1/4} \rfloor^4 + 1} + \dots + X_n) / n = Y_{\lfloor n^{1/4} \rfloor^4} + (M + |\mu|) \frac{n - \lfloor n^{1/4} \rfloor^4}{n}$.
The first term converges to μ as $n \rightarrow \infty$, and the second converges to 0.

3 Sample statistics

3.1 Some important distributions

- Standard Normal Distribution: $\mathcal{N}(0, 1)$
- $\chi^2(k)$: squared sum of k independent standard normal distribution.
- t distribution: Z standard normal, $Y \sim \chi^2(k)$, Z and Y independent, then $T = \frac{Z}{\sqrt{Y/k}}$ is said to have t -distribution with k degrees of freedom.
- F distribution: U and V independent, $U \sim \chi^2(m)$, $V \sim \chi^2(n)$, then $F = \frac{U/m}{V/n}$ is said to have F distribution with degrees of freedom m and n ,

3.2 Sample statistics

X_1, \dots, X_n i.i.d. (independent with identical distributions). Sample statistics: a random variable computed from n other random variables.

- **Sample mean:** $\bar{X} = \frac{\sum_i X_i}{n}$

$$- E[\bar{X}] = E[X_1], \text{Var}(\bar{X}) = \frac{1}{n} \text{Var}(X_1).$$

Proof:

$$E[\bar{X}] = E\left[\frac{1}{n} \sum_i X_i\right] = \frac{1}{n} \sum_i E[X_i] = E[X_1]$$

$$\text{Var}(\bar{X}) = E[(\bar{X} - E[X_1])^2] = \frac{1}{n^2} E\left[\sum_i (X_i - E[X_i])^2\right] = \frac{1}{n} \text{Var}(X_1)$$

- If $X_1 \sim \mathcal{N}(\mu, \sigma^2)$, $\bar{X} \sim \mathcal{N}(\mu, \sigma^2/n)$.

Proof: By calculation using MGF.

- If $n \rightarrow \infty$, $\sqrt{\frac{n}{\text{Var}(X_1)}}(\bar{X} - E[X_1])$ converges to standard normal by distribution.

Proof: This is just central limit theorem.

- **Sample variance:** $S^2 = \frac{1}{n-1} \sum_i (X_i - \bar{X})^2 = \frac{1}{n-1} (\sum_i X_i^2 - n\bar{X}^2)$.

$$- E[S^2] = \text{Var}(X_1).$$

Proof:

$$E[S^2] = \frac{1}{n-1} \sum_i E[(X_i - \bar{X})^2] = \frac{1}{n-1} \sum_i E\left[\left(\frac{n-1}{n} X_i - \sum_{j \neq i} \frac{1}{n} X_j\right)^2\right]$$

$$\begin{aligned}
&= \frac{1}{n-1} \sum_i \left(\frac{(n-1)^2}{n^2} E[X_i^2] + \sum_{j \neq i} \frac{1}{n^2} E[X_j^2] - \sum_{j \neq i} \frac{2n-2}{n^2} E[X_i] E[X_j] \right. \\
&\quad \left. + \sum_{j \neq i, k \neq i, j \neq k} \frac{2}{n^2} E[X_j] E[X_k] \right) \\
&= E[X_1^2] - E[X_1]^2 = \text{Var}(X_1)
\end{aligned}$$

– If $X_1 \sim \mathcal{N}(\mu, \sigma^2)$, then

* \bar{X} and S^2 are independent

Proof: Calculate joint cdf, do a change of variables.

* $\frac{(n-1)S^2}{\sigma^2} \sim \chi^2(n-1)$

Proof:

$$\frac{(n-1)S^2}{\sigma^2} + n \frac{(\bar{X} - E[X_1])^2}{\sigma^2} = \frac{1}{\sigma^2} \sum_i (X_i - E[X_1])^2 \sim \chi^2(n)$$

Now use moment generating function and the independence between S^2 and \bar{X} .

* $\frac{\bar{X} - \mu}{S/\sqrt{n}} \sim t(n-1)$.

Proof: By definition of t -distribution.

– If S_1^2 is the sample variance of n_1 i.i.d. $\mathcal{N}(\mu, \sigma^2)$ random variables Y_i , S_2^2 the sample variance of n_2 i.i.d. $\mathcal{N}(\mu', \sigma'^2)$ random variables Z_j independent from Y_i , then $\frac{S_1^2/\sigma_1^2}{S_2^2/\sigma_2^2} \sim F(n_1-1, n_2-1)$

Proof: By definition of F -distribution.

- **Order statistics** The k -th order statistics is the k -th smallest element in $\{X_i\}$, denoted as Y_k . Then, if X_1 has pdf f , then

$$\begin{aligned}
f_{Y_k}(t) &= \frac{d}{dt} F_{Y_k}(t) = \lim_{\delta \rightarrow 0} \frac{F_{Y_k}(t+\delta) - F_{Y_k}(t)}{\delta} \\
&= \lim_{\delta \rightarrow 0} \frac{1}{\delta} \binom{n}{k-1, 1, n-k} \left(\int_{-\infty}^t f ds \right)^{k-1} \int_t^{t+\delta} f ds \left(\int_{t+\delta}^{\infty} f ds \right)^{n-k} \\
&= \frac{n!}{(k-1)!(n-k)!} \left(\int_{-\infty}^t f ds \right)^{k-1} f(t) \left(\int_t^{\infty} f ds \right)^{n-k}
\end{aligned}$$

3.3 PDF of χ^2 -, t- and F- distributions

3.3.1 χ^2

Let X_i be iid standard normal, their joint distribution is

$$f(x_1, \dots, x_n) = (2\pi)^{-n/2} e^{-\sum_i x_i^2/2}$$

Hence the pdf of χ^2 is:

$$f_{\chi^2(n)}(r) = \frac{d}{dr} \int_{\sum_i x_i^2 \leq r} (2\pi)^{-n/2} e^{-\sum_i x_i^2/2} dx_1 \dots dx_n$$

which is easy to see must be proportional to $r^{\frac{n-2}{2}} e^{-r/2}$.

3.4 t

Let X and Y be independent with pdf: $f_X(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}$ and $f_Y(y) = \frac{1}{2^{d/2}\Gamma(d/2)} y^{\frac{d-2}{2}} e^{-y/2}$. Then

$$\begin{aligned} f_{t(d)}(s) &= \frac{d}{ds} P(X \leq s\sqrt{Y/d}) = \frac{d}{ds} \int_0^\infty dy \int_{-\infty}^{s\sqrt{y/d}} dx \frac{1}{\sqrt{2\pi}} e^{-x^2/2} \frac{1}{2^{d/2}\Gamma(d/2)} y^{\frac{d-2}{2}} e^{-y/2} \\ &= \int_0^\infty dy \sqrt{y/d} \frac{1}{\sqrt{2\pi}} e^{-s^2 y/2d} \frac{1}{2^{d/2}\Gamma(d/2)} y^{\frac{d-2}{2}} e^{-y/2} \end{aligned}$$

Do change of variables $z = (s^2/d + 1)y$ we get that it is proportional to $(s^2/d + 1)^{-\frac{d+1}{2}}$.

The calculation for the pdf of F is similar.

4 Point estimators and their properties

Basic setting:

- \mathcal{F} : a family of possible distributions (represented by a family of cdf, pdf, or pd)
- $\theta : \mathcal{F} \rightarrow \mathbb{R}$ population parameter
- X_1, \dots, X_n i.i.d. with distribution $F \in \mathcal{F}$
- $\hat{\theta} = \hat{\theta}(X_1, \dots, X_n)$ a function of X_i , which is an estimate of $\theta(F)$, is called a point estimate.

Example: \mathcal{F} : all distributions with an expectation, then \bar{X} is a point estimate of the expectation.

$\hat{\theta}$ is a point estimate of θ .

- The **bias** is $E[\hat{\theta}] - \theta$. $\hat{\theta}$ is called unbiased if $E[\hat{\theta}] = \theta$.
- The **variance** is $Var(\hat{\theta})$.
- $\hat{\theta}$ is called **minimum variance unbiased estimate** if it has the smallest variance among all unbiased estimates.
- $\hat{\theta}_1$ and $\hat{\theta}_2$ are two unbiased estimates, the relative efficiency is the ratio of their variance. When they are biased, one can use the mean squared error $E[(\hat{\theta} - \theta)^2]$ instead.
- $\hat{\theta}$ is called **asymptotically unbiased** if bias converges to 0 as $n \rightarrow \infty$.
- $\hat{\theta}$ is called **consistent** if $\hat{\theta}$ converges to θ in distribution.

Example: Estimate of the expectation and variance of binomial distribution

- Expectation can be estimated by sample mean, which is unbiased and consistent.
- Variance can be estimated by sample variance which is unbiased and consistent, or $\bar{X}(1 - \bar{X})$, which is consistent but biased.

Example: Estimate t for uniform distribution on $[0, t]$.

The following estimates are all unbiased and consistent:

- $2\bar{X}$
- $\frac{n+1}{n} \text{Max}(X_i)$
- $\text{Max}(X_i) + \text{Min}(X_i)$

Can you calculate their variance? Which is the best among the three?

Answer:

$$\begin{aligned}
 Var(2\bar{X}) &= \frac{4}{n} \cdot Var(X_1) = \frac{t^2}{3n} \\
 Var\left(\frac{n+1}{n} \text{Max}(X_i)\right) &= \frac{(n+1)^2}{n^2} \cdot n! \cdot \int_0^t dx_n \int_0^{x_n} dx_{n-1} \cdots \int_0^{x_2} dx_1 \cdot \frac{(x_n - t)^2}{t^n} \\
 &= \frac{(n+1)^2}{n} \int_0^t \frac{(x_n - \frac{nt}{n+1})^2 x_n^{n-1}}{t^n} dx_n = \frac{t^2}{n(n+2)} \\
 Var(\text{Max}(X_i) + \text{Min}(X_i)) &= \frac{n!}{t^n} \cdot \int_0^t dx_n \int_0^{x_n} dx_1 \int_{x_1}^{x_n} dx_{n-1} \cdots dx_2 \cdot (x_n + x_1 - t)^2 \\
 &= \frac{n(n-1)}{t^n} \int_0^t dx_n \int_0^{x_n} dx_1 (x_n + x_1 - t)^2 (x_n - x_1)^{n-2} = \frac{2t^2}{(n+1)(n+2)}
 \end{aligned}$$

If an asymptotically unbiased estimate has variance $\rightarrow 0$ when $n \rightarrow \infty$, it must be consistent.

Cramer-Rao inequality:

$$\text{Var}(\hat{\theta}) \geq \frac{1}{nE[(\frac{d}{d\theta} \log f)^2]}$$

When equality is reached we get minimal variance unbiased estimate.

Example: X_i iid normal, then \bar{X} is MVUE.

$$\text{Var}(\bar{X}) = \sigma^2/n$$

$$\frac{1}{nE[(\frac{d}{d\theta} \log f)^2]} = \frac{1}{nE[(X - \mu)^2/\sigma^4]} = \sigma^2/n$$

5 Method of moments, Maximum likelihood

5.1 MLE

Suppose $X_i \sim F \in \mathcal{F}$, i.i.d., where \mathcal{F} is the family of possible distributions of X_i , and F is unknown and belongs to \mathcal{F} . We want to find a point estimate for some function $\theta : \mathcal{F} \rightarrow \mathbb{R}$. The Method of Maximal Likelihood is:

$$\hat{\theta}(X_1, \dots, X_k)_{MLE} = \theta(\arg \max_{F \in \mathcal{F}} L(X_1, \dots, X_k, F))$$

- When F is a continuous distribution with p.d.f. $f(x)$, let $L(x_1, \dots, x_k, F) = \prod_i f(x_i)$
- When F is a discrete distribution with p.d. $g(x) = P(X = x)$, let $L(x_1, \dots, x_k, F) = \prod_i g(x_i)$

Example: X_i i.i.d. and has binomial distribution with $n = 5$ and unknown p , find MLE for p .

Answer: If X_i satisfies the binomial distribution with $n = 5$ and let p be some unknown value, the likelihood function is:

$$L(X_1, \dots, X_k) = \prod_i \binom{5}{X_i} p^{X_i} (1-p)^{5-X_i}$$

The p that maximizes it is $p = \frac{\sum_i X_i}{5k}$, hence $\hat{p}_{MLE} = \frac{\sum_i X_i}{5k}$

Example: X_i i.i.d. and has uniform distribution on $[a, a + t]$. Find MLE for a and t .

Answer: If $[a, a + t]$ fails to contain any of the X_i the likelihood must be 0, so $a \leq \min\{X_i\}$, $a + t \geq \max\{X_i\}$. To maximize the likelihood in this case, one need to minimize t , hence $\hat{a}_{MLE} = \min\{X_i\}$ and $\hat{t}_{MLE} = \max\{X_i\} - \min\{X_i\}$.

Example: X_i i.i.d. and has normal distribution with expectation μ variance σ^2 . Find MLE for σ^2 .

Answer: Write down the likelihood function, take derivative for both μ and σ^2 and set both to be 0, we get that $\hat{\sigma}^2_{MLE} = \frac{1}{n} \sum_i (X_i - \bar{X})^2$.

5.2 MOM

MOM is a less popular approach but does have some advantages in some situations.

Empirical distribution: Given $x_1, \dots, x_k \in \mathbb{R}$, the empirical distribution X' is defined as $P(X' = x_i) = \frac{m_i}{k}$ where m_i is the multiplicity of x_i .

Method of moments means estimating the parameters in such a way that the first few moments of X_i under these parameters match the first few moments of empirical distribution obtained from X_1, \dots, X_k , i.e. the sample moments $M'_n = \frac{1}{k} \sum_i X_i^n$.

Example: X_i i.i.d. uniform on $[a, a + t]$, find MOM estimate for a and t .

Example: X_i i.i.d. exponential, $f(x) = \frac{1}{c} e^{-x/c}$, find MOM estimate for c .

Example: X_i i.i.d. binomial with $p = \frac{1}{2}$. Find MOM and MLE for n . Are they the same?

6 Point estimate for non i.i.d. random variables

- \mathcal{F} : a family of possible joint distributions (represented by a family of joint cdf, joint pdf, or joint pd)
- $\theta : \mathcal{F} \rightarrow \mathbb{R}$ population parameter
- $X_1, \dots, X_n \sim F \in \mathcal{F}$
- $\hat{\theta} = \hat{\theta}(X_1, \dots, X_n)$ a function of X_i , which is an estimate of $\theta(F)$, is called a point estimate.

One can define bias, variance and consistency similar to the i.i.d. case. The MLE (and MAP which will be discussed later) works for non i.i.d. case as well!

Example: X_1, \dots, X_n uniform on $[a, a + t]$, Y_1, \dots, Y_n uniform on $[b, b + t]$, find MLE of t .

7 Midterm 1 review

Some key topics:

- pdf from cdf, probability from pdf, expectation from pdf.
- LLN and CLT
- Sample mean and sample variance: general case and normal population
- Bias, variance, mean squared error and consistency for point estimate
Ways to check consistency:
 - Definition
 - LLN
 - mean squared error
- MLE

Review problem:

1. X_i i.i.d. $P(X_i = 0) = a$, $P(X_i = 1) = b$, $P(X_i = 2) = c$, $a + b + c = 1$. Find $E[X_i]$, the MLE of $E[X_i]$, and the bias and variance of said MLE. Find $Var(X_i)$ and its MLE. Are these MLEs consistent?

2. X_i i.i.d., with pdf $\frac{1}{\sqrt{2\pi}}(ce^{-x^2/2} + (1 - c)e^{-(x-1)^2/2})$. Find MLE of c . Is it consistent?

8 Digression: The idea of substitution

Examples:

- If X_i i.i.d. with p.d.f. f , then likelihood function $L = \prod_i f(X_i)$. If Y_i i.i.d., tY_i has $\chi^2(2)$ distribution (p.d.f. $f(x) = \frac{1}{2}e^{-x/2}$ when $x \geq 0$), what is the likelihood function?
- Let $f_\theta(\cdot)$ be the p.d.f. of θ , $f_{X|\theta}(\cdot, \theta)$ the conditional p.d.f. of X , $f_{\theta|X}(\cdot, X)$ the conditional p.d.f. of θ , then

$$f_{\theta|X}(\theta, x) = \frac{f_\theta(\theta)f_{X|\theta}(x, \theta)}{\int_{\mathbb{R}} f_\theta(\theta)f_{X|\theta}(x, \theta)d\theta}$$

9 Bayesian statistics

9.1 The basic idea of Bayesian statistics

- Input:
 - Some (possibly vector valued) random variable Θ with given distribution (**prior**)
 - Some (possibly vector valued) random variable X with known conditional distribution conditioned at a value of Θ , $X \sim F(X|\Theta)$. (**observable**)
- Output: the conditional distribution of Θ conditioned at a value of X (**posterior**) $\Theta \sim F(\Theta|X)$.

Example 1:

- **Prior** $Y \sim \text{Bernoulli}(\frac{1}{100})$
- **Observable** X_1, X_2 conditionally i.i.d. when $Y = y$, and their conditional distribution is Bernoulli with $p = \frac{1+8Y}{10}$.

Calculation of the posterior:

$$\begin{aligned}
 P(Y = 1|X_1, X_2) &= \frac{P(Y = 1, X_1, X_2)}{P(X_1, X_2)} \\
 &= \frac{P(X_1, X_2|Y = 1)P(Y = 1)}{P(X_1, X_2|Y = 0)P(Y = 0) + P(X_1, X_2|Y = 1)P(Y = 1)} \\
 &= \frac{(9/10)^{X_1+X_2}(1/10)^{2-X_1-X_2} \times \frac{1}{100}}{(9/10)^{X_1+X_2}(1/10)^{2-X_1-X_2} \times \frac{1}{100} + (1/10)^{X_1+X_2}(9/10)^{2-X_1-X_2} \times \frac{99}{100}} \\
 &= \frac{9^{X_1+X_2}}{9^{X_1+X_2} + 99 \times 9^{2-X_1-X_2}}
 \end{aligned}$$

So, for example, if we know both X_i takes a value of 1, then the probability of $Y = 1$ is $9/20$.

We can answer many questions using posterior, for example:

- What is the probability of Θ taking value in A given X ?
- What is the “most likely” value of Θ ? $\hat{\Theta}_{MAP} = \arg \max_s f_{\Theta|X}(s)$, where f is p.d.f. when $\Theta|X$ is continuous and p.d. when it is discrete. This is called the **maximum a posteriori (MAP)** estimate.
- What is the average value of Θ ? $\hat{\Theta} = E[\Theta|X]$. This is called the **Bayesian point estimate with L^2 lost**.
- In general, let $l(\cdot, \cdot)$ be a lost function (a positive function such that $l(a, a) = 0$), then $\hat{\Theta} = \arg \min_{\theta} E[l(\Theta, \theta)|X]$ is called the **Bayesian point estimate**.

9.2 Comparison between non-Bayesian and Bayesian

MLE:

- Input: Assumption on the distribution of X : $X \sim F(\alpha)$. A likelihood function $L(X, \alpha)$.
- Output: $\hat{\alpha}_{MLE} = \arg \max_{\alpha} L(X, \alpha)$.

Bayesian statistics:

- Input: Prior: $\alpha \sim F_0$, Conditional distribution: $X|\alpha \sim F(\alpha)$.
- Calculated output: Posterior: $\alpha|X \sim F'(X)$
- MAP Point estimate: $\hat{\alpha} = \arg \max_{\alpha} f_{\alpha|X}(\alpha)$
- L^2 -Bayesian Point estimate: $\hat{\alpha} = E[\alpha|X]$.

9.3 Example of point estimate using Bayesian statistics

Example 2:

Input:

- $\mu \sim \mathcal{N}(0, 1)$
- $X_i|\mu$ cond. i.i.d., $\sim \mathcal{N}(\mu, 1)$

Posterior:

$$\begin{aligned} f_{\mu|X_i}(s) &= \frac{f_{\mu, X_i}(s, X_1, \dots, X_n)}{f_{X_i}(X_1, \dots, X_n)} = \frac{f_{\mu, X_i}(s, X_1, \dots, X_n)}{\int_{\mathbb{R}} f_{\mu, X_i}(t, X_1, \dots, X_n) dt} \\ &= \frac{\prod_i f_{X_i|\mu=s}(X_i) f_{\mu}(s)}{\int_{\mathbb{R}} \prod_i f_{X_i|\mu=t}(X_i) f_{\mu}(t) dt} = \frac{(2\pi)^{-\frac{n+1}{2}} e^{-\sum_i (X_i-s)^2/2 - s^2/2}}{\int_{\mathbb{R}} (2\pi)^{-\frac{n+1}{2}} e^{-\sum_i (X_i-t)^2/2 - t^2/2} dt} \end{aligned}$$

So

$$\mu|X_i \sim \mathcal{N}\left(\frac{\sum_i X_i}{n+1}, \frac{1}{n+1}\right)$$

The MAP and L^2 Bayesian estimate of μ are both $\hat{\mu} = \frac{\sum_i X_i}{n+1}$.

From the computation above we get:

$$f_{\mu|X}(s) \propto f_{X|\mu=s}(X) f_{\mu}(s)$$

This works for discrete μ or X as well!

Example 3: P uniform on $[0, 1]$, $X|P \sim \text{Binomial}(5, P)$, then $f_{P|X}(s) \propto s^X (1-s)^{5-X} \cdot 1$, hence $P|X \sim \text{Beta}(X+1, 6-X)$.

9.4 Hierarchical Models

This section is beyond the scope of our exams.

Often in practice we build “hierarchical models” by stacking multiple layers of Bayesian and non Bayesian models together. For example:

$$\begin{aligned}\sigma_i^2 &\sim \Gamma(\alpha, \beta) \\ \sigma^2 &\sim \Gamma(\alpha', \beta') \\ \mu_i &\sim \mathcal{N}(0, \sigma^2) \\ X_{ij} \text{ ind. } &\sim \mathcal{N}(\mu_i, \sigma_i^2)\end{aligned}$$

How would you estimate σ_i and μ_i from the values of X_{ij} ?

We will talk about models like this if we have more time at the end of the semester.

9.5 More Examples

Example 4: t has p.d.f. $f_t(x) = \begin{cases} 0 & x < 0 \\ e^{-x} & x > 0 \end{cases}$. $P(Y = n|t) = (1 - e^{-t})e^{-nt}$.
Knowing Y , find \hat{t}_{MAP} and $E[t|Y]$.

Example 5: a, t indep. $\sim \text{Uniform}([0, 1])$. $X_i|a, t$ i.i.d. $\sim \text{Uniform}([a, a + t])$, find \hat{t}_{MAP} .

Answer: $M = \max(X_i)$, $m = \min(X_i)$, then:

$$f_{a,t|X_i} \propto \begin{cases} t^{-n} & 0 \leq a \leq m \leq M \leq a + t \leq a + 1 \\ 0 & \text{otherwise} \end{cases}$$

So

$$\begin{aligned}f_{t|X_i} &\propto \begin{cases} t^{-n} \cdot (\min(1, m) - (M - t)) & M - \min(1, m) \leq t \leq 1 \\ 0 & \text{otherwise} \end{cases} \\ \hat{t}_{MAP} &= \min(1, \frac{n}{n-1}(M - \min(1, m)))\end{aligned}$$

10 Hypothesis testing

10.1 Definitions

- Problem: want to know if the distribution of X satisfy certain propositions (**null hypothesis**), for example:
 - Will the coronavirus kill more than a million people in the end?

- Will the expectation of our midterm 2 grade be better than midterm 1?
- Is the performance of a machine learning algorithm better than random chance?
- Solution: Find a random variable Z (**test statistics**) depending on X and a set A (**critical region**), and reject the hypothesis when $Z \in A$.
- (Z, A) is called a **statistical test** to null hypothesis H_0 .
- If $Z \in A \iff Z' \in A'$ we consider (Z, A) and (Z', A') to be the same test.
- If H_0 completely determines $P(Z \in A)$ (**simple hypothesis**), $p = P(Z \in A|H_0)$ is called the **p-value**.

The key reasoning behind statistical tests: **Suppose H_0 is true. If (Z, A) is a test with a very small p -value, then $Z \in A$ is highly unlikely. If, however, we do actually observe that $Z \in A$, then this can only tell us that H_0 is unlikely to be true.** It is basically a kind of “statistical” proof by contradiction.

Example 1: Suppose your grade for midterm 1 is X_1 , your grade for midterm 2 is X_2 , $Y = X_2 - X_1$ satisfies normal distribution with variance 25. How do we test the null hypothesis $E[Y] = 0$?

- Answer 1: $Z = Y$, $A = (-\infty, -M) \cup (M, \infty)$.

$$\begin{aligned} p &= P(Y < -M \cup Y > M | H_0) = P(Y < -M | Y \sim \mathcal{N}(0, 25)) + P(Y > M | Y \sim \mathcal{N}(0, 25)) \\ &= 2 \int_M^\infty \frac{1}{\sqrt{50\pi}} e^{-t^2/50} dt \end{aligned}$$

- Answer 2: $Z = Y$, $A = (M, \infty)$, $p = \int_M^\infty \frac{1}{\sqrt{50\pi}} e^{-t^2/50} dt$
- Answer 3: $Z = Y$, $A = (-M, M)$, $p = \int_{-M}^M \frac{1}{\sqrt{50\pi}} e^{-t^2/50} dt$

Which of the three is more reasonable?

- **Alternative hypothesis:** an alternative to the null hypothesis H_0 , called H_1 .
- $P(Z \in A | H_0)$ is called **p-value** or **type I error**.
- If H_1 is a simple hypothesis, $P(Z \notin A | H_1)$ is called **type II error**.
- If H_1 is a simple hypothesis, $1 - P(Z \notin A | H_1) = P(Z \in A | H_1)$ is called **(statistical) power**

- If $X \sim F(\theta)$, $\pi(\theta) = P(Z \in A|\theta)$ is called the **power function**. If $H_0 : \theta = \theta_0$, $H_1 : \theta = \theta_1$, then p-value is $\pi(\theta_0)$ and power is $\pi(\theta_1)$.

In Example 1, let $Y = \mathcal{N}(\theta, 25)$, what is the power function of the three tests?

Answer: Let $f(y) = \frac{1}{\sqrt{50\pi}} e^{-y^2/50}$. Then for Test 1,

$$\pi_1(\theta) = \int_{(-\infty, -M) \cup (M, \infty)} f(t - \theta) dt = \int_{(-\infty, -M - \theta) \cup (M - \theta, \infty)} f(s) ds$$

So $\frac{d\pi}{d\theta} = f(M - \theta) - f(-M - \theta)$, which is positive when $\theta > 0$ and negative when $\theta < 0$. So, if the alternative hypothesis is $\theta = \theta_1 > 0$ or $\theta = \theta_1 < 0$, it is possible to find some M which make p-value small and power large.

For Test 2,

$$\pi_2(\theta) = \int_{(M, \infty)} f(t - \theta) dt = \int_{(M - \theta, \infty)} f(s) ds$$

So $\frac{d\pi}{d\theta} = f(M - \theta) > 0$. So if the alternative hypothesis is $\theta = \theta_1 > 0$ it is possible to find some M which make p-value small and power large. In other words, this test can only capture the case when $E[Y] > 0$ but not $E[Y] < 0$, which is consistent with our expectation.

For Test 3, the power function is

$$\pi_3(\theta) = \int_{(-M, M)} f(t - \theta) dt = \int_{(-M - \theta, M - \theta)} f(s) ds$$

So $\frac{d\pi}{d\theta} = -f(M - \theta) + f(-M - \theta)$ which is negative when $\theta > 0$ and positive when $\theta < 0$, so as a consequence the type I and type II errors always sum up to something larger than 1, which means that it is a very bad test.

Example 2: Y_i i.i.d. $\sim \mathcal{N}(\theta, 25)$, $H_0 : \theta = 0$. What is the power function for the test $(\bar{Y}, (-\infty, -M) \cup (M, \infty))$?

Example 3: Y_i i.i.d. Bernoulli distribution with parameter θ , $H_0 : \theta = 0.5$. What is the power function for the test $(\bar{Y}, (0, 1/2 - \epsilon) \cup (1/2 + \epsilon, 1))$?

Example 4: X_i $i = 1, \dots, 6$ i.i.d., Bernoulli with $P(X_i = 1) = p$. $H_0 : p = 0.5$, $H_1 : p = 0.9$. Test statistics: $Z = \sum_i X_i$. $A = [M, 6]$, M is an integer.

Then power function is:

$$\pi(p) = P(Z \geq M|p) = \sum_{i=M}^6 \binom{6}{i} p^i (1-p)^{6-i}$$

p-value is $\pi(0.5) = \frac{1}{64} \sum_{i=M}^6 \binom{6}{i}$. Power is $\pi(0.9) = \sum_{i=M}^6 \binom{6}{i} (0.9)^i (0.1)^{6-i}$.

- $M = 6$: p-value=0.0156, power=0.531
- $M = 5$: p-value=0.109, power=0.886
- $M = 4$: p-value=0.344, power=0.984

There is trade-off between p-value and power. Which M to choose depends on the purpose of the test, in particular whether false positive or false negative would be more costly.

10.2 Likelihood ratio test

Recall that the likelihood function is $L(x, \theta) = f_{X|\theta}(x)$, which is the p.d.f. when X is continuous and p.d. when X is discrete. The Neyman-Pearson test for $H_0 : \theta = \theta_0$, $H_1 : \theta = \theta_1$ is:

$$(X, \{x : L(x, \theta_0)/L(x, \theta_1) \leq k\})$$

Example 4, Neyman-Pearson test: $p_0 = 0.5$, $p_1 = 0.9$

$$L(X_1, \dots, X_6, p_0) = \prod_i p_0^{X_i} (1 - p_0)^{1-X_i} = \frac{1}{2^6}$$

$$\begin{aligned} L(X_1, \dots, X_6, p_1) &= \prod_i p_1^{X_i} (1 - p_1)^{1-X_i} \\ &= 0.9^{\sum_i X_i} \cdot 0.1^{6 - \sum_i X_i} = 0.1^6 \cdot 9^{\sum_i X_i} \end{aligned}$$

So likelihood ratio decreases with $\sum_i X_i$.

Sometimes we need to consider **composite hypothesis**, i.e. cases when H_0 and H_1 does not completely determine the distribution of X . Suppose $H_0 : \theta \in D_0$, $H_1 : \theta \in D_1$, the likelihood ratio test becomes:

$$(X, \{x : \frac{\sup_{\theta \in D_0} L(x, \theta)}{\sup_{\theta \in D_0 \cup D_1} L(x, \theta)} \leq k\})$$

How would you do likelihood ratio test for the following examples:

- X_i i.i.d. Bernoulli(p). $H_0 : p = 0.5$, $H_1 : p \neq 0.5$.
- X_i i.i.d. $\mathcal{N}(\mu, 1)$. $H_0 : \mu = 0$, $H_1 : \mu \neq 0$.

Answer:

- Likelihood under H_0 is

$$L_0 = \prod_i 0.5^{X_i} (1 - 0.5)^{1-X_i} = 0.5^n$$

maximum likelihood under H_1 is

$$\begin{aligned} L_1 &= \sup_p \prod_i p_i^X (1-p)^{1-X_i} \\ &= \sup_p p^{\sum_i X_i} (1-p)^{n-\sum_i X_i} \\ &= (\sum_i X_i/n)^{\sum_i X_i} (1-\sum_i X_i/n)^{n-\sum_i X_i} \end{aligned}$$

It is easy to see that the likelihood ration L_0/L_1 , as a function of $\sum_i X_i$, is symmetric with regards to $n/2$, and takes its maximum at $\sum_i X_i = n/2$. So the likelihood ratio test must be of the form: $|\sum_i X_i - n/2| \geq C$ for some C .

- Likelihood under H_0 is

$$L_0 = \prod_i \frac{1}{\sqrt{2\pi}} e^{-X_i^2/2} = (2\pi)^{-n/2} e^{-\frac{\sum_i X_i^2}{2}}$$

maximum likelihood under H_1 is

$$\begin{aligned} L_1 &= \sup_\mu \prod_i \frac{1}{\sqrt{2\pi}} e^{-(X_i-\mu)^2/2} \\ &= \sup_\mu (2\pi)^{-n/2} e^{-\frac{\sum_i (X_i-\mu)^2}{2}} \\ &= (2\pi)^{-n/2} e^{-\frac{\sum_i X_i^2 - (\sum_i X_i)^2/n}{2}} \end{aligned}$$

So

$$L_0/L_1 = e^{-\frac{(\sum_i X_i)^2}{2n}}$$

So the likelihood ratio test must be of the form $|\sum_i X_i| \geq C$.

- 11 Examples of hypothesis testing**
- 12 Confidence interval**
- 13 Linear Regression**
- 14 ANOVA**
- 15 Example of non parametric methods**